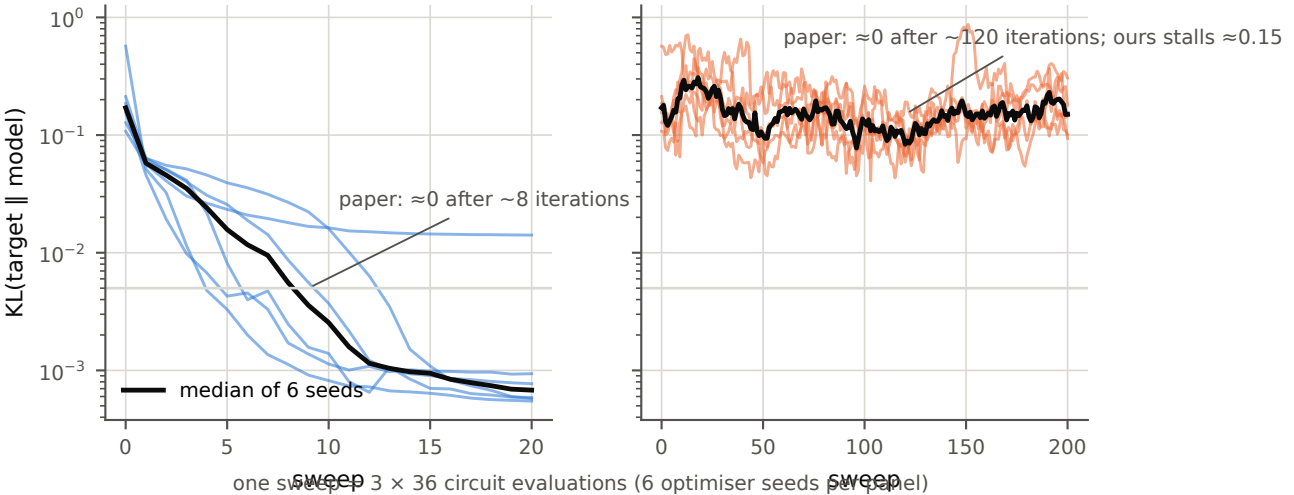
superposition training
iterative (one input per step) — our construction
KL(target ‖ model)
10^{0}
10^{-1}
10^{-2}
10^{-3}
paper: ≈0 after ~8 iterations
median of 6 seeds
paper: ≈0 after ~120 iterations; ours stalls ≈0.15
0
5
10
15
20
50
100
150
200
sweep
sweep
one sweep = 3 × 36 circuit evaluations (6 optimiser seeds per panel)