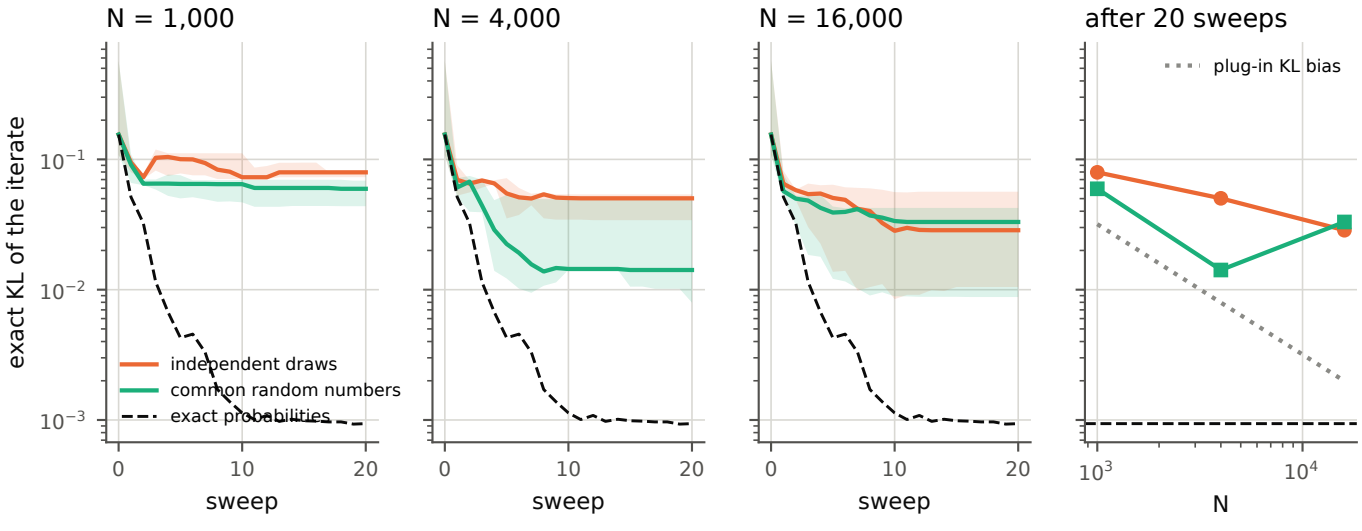
Bit-count KW–SA training when every loss evaluation is estimated from N shots (3 seeds: median, min–max)
N = 1,000
N = 4,000
N = 16,000
after 20 sweeps
exact KL of the iterate
$10^{-1}$
$10^{-2}$
$10^{-3}$
independent draws
common random numbers
exact probabilities
plug-in KL bias
0
10
20
sweep
$10^{3}$
$10^{4}$
N