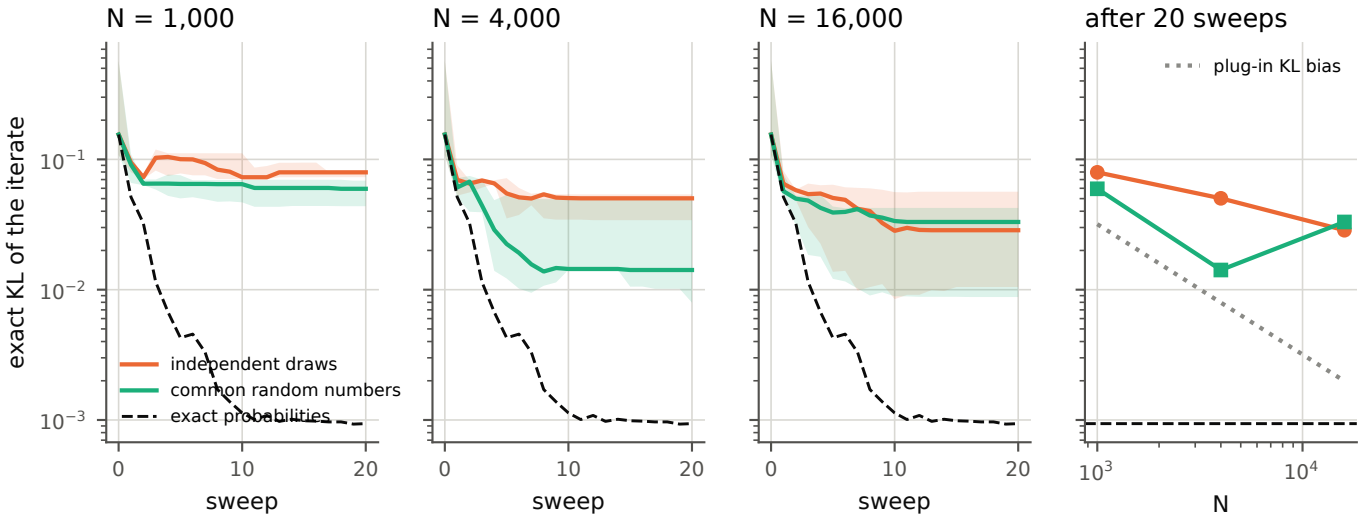
Bit-count KW–SA training when every loss evaluation is estimated from N shots (3 seeds: median, min–max)
N = 1,000
N = 4,000
N = 16,000
after 20 sweeps
exact KL of the iterate
10^{-1}
10^{-2}
10^{-3}
independent draws
common random numbers
exact probabilities
plug-in KL bias
0
10
20
sweep
10^{3}
10^{4}
N